

FREE STARTER GUIDE

The 4-Folder System

How AI goes from slop to sounding like you

A 14-page primer extracted from the AI Mastery Practice Manual.
Read it. Try the exercise on page 6. Then decide if the full course is for you.

Welcome — and a promise

If your AI output sounds generic, it is not the model failing you. It is missing three things — and the rest of this guide hands them to you.

By the time you finish reading (about 20 minutes), you will:

- Understand exactly why ChatGPT, Claude, and Gemini produce slop by default.
- Have a copy-paste folder structure you can set up in 5 minutes.
- Know what to put in the single most important file in your AI workflow.
- Know exactly what to upgrade to if you want the rest of the system.

KEY IDEA

This is a real chapter from the full Practice Manual — not a watered-down teaser. Three chapters in, you will know whether the rest of the course is for you.

Chapter 1 — The slop problem

Why your AI output sounds generic

Most people use AI by typing a one-line request and pasting the result. The output sounds vaguely correct, vaguely on-brand, and entirely forgettable. The internet has a name for this: AI slop.

Slop is not a model problem. It is an input problem. The model has nothing to anchor on except its training data — the average of the internet — so it produces the average of the internet back.

KEY IDEA

If your AI output sounds generic, the model is not failing. You are.

The fix is upstream: who is asking, what they have done before, and what they are asking for.

Signal in, signal out

Three inputs decide the quality of every AI run:

1. Business context — the who. Who is asking, what they sell, the world they inhabit.
2. Data — the past. The author's actual prior work, which encodes voice, structure, and taste.
3. Prompt — the ask. The shape of the request: role, process, constraints, format.

Skip any one and you get slop. Provide all three and the same model that wrote the slop now writes in your voice.

Proof: one idea, three runs, same model

Below is the same idea — "write a tweet about pricing mistakes" — sent to the same model three times. The only thing that changes is what we attached.

Run A — prompt only

Prompt:

Write a tweet about common pricing mistakes founders make.

Result:

"Many founders underprice. Don't leave money on the table.
Charge what you're worth! #startup #pricing"

Generic. Hashtags. Could have been written by anyone.

Run B — prompt + business context

Attached: business_context.md (1 page about a B2B agency owner).

Result:

"Charging hourly turned my agency into a busy job.
The day we switched to monthly retainers, profit per hour 2.5x.
Same work. Different invoice."

Better. The model now knows the speaker is an agency owner; the example is concrete and on-brand.
Still missing the speaker's actual sentence rhythm.

Run C — prompt + business context + 200 of the author's past tweets

Result:

"Hourly rates put a ceiling on you.
Retainers put a floor under you.
Same hours. Different ladder."

Now the rhythm matches. Two short setups, one short payoff. No hashtags. The author's voice.

*"AI sounds like the internet because it is trained on the internet. Train it on your own
work, and it sounds like your own work."*

— the core principle

Chapter 2 — The 4-folder system

Once you accept that quality output requires three inputs, the obvious question is: where do they live? The system you're about to set up takes 5 minutes and never changes.

The structure

```
AI_Mastery/  
  01_Business_Context/    ← single source of truth about you/your company  
  02_Data/                ← your historical work, by content type  
    tweets/  
    youtube_transcripts/  
    shortform_transcripts/  
  03_Prompts/             ← reusable prompt templates per task  
  04_Code_Projects/       ← (we cover this in Module 3 of the full course)
```

Why this exact shape

- **Single source of truth.** Business context lives in exactly one file. Every prompt references it. When your offer changes, you edit one place.
- **Voice fidelity.** Data is organized by output type because voice differs by format. Your tweet voice is not your YouTube voice.
- **Prompt portability.** Prompts live separately from data so you can re-use the same prompt across projects with different data.

Set it up in 5 minutes

1. Create the four folders on your computer. Name them exactly as above (the numeric prefix keeps them sorted).
2. Move any business documents you already have into 01_Business_Context/.
3. Dump your latest 50 social posts into 02_Data/tweets/ as plain .txt files — one per file.
4. Open a new chat with Claude. Drag the entire AI_Mastery folder in. The model now has the entire structure as context.

WIN PATTERN

Treat 01_Business_Context/ like the README of your professional life. If it goes stale, your AI outputs will too. Re-read it monthly.

Chapter 3 — Your business context document

This is the single most important asset in the entire system. The document is the model's only window into who you are. A weak business context creates weak everything downstream.

KEY IDEA

Spend an afternoon on this. Most people spend 10 minutes and then complain that the AI sounds generic.

The 12-question framework

Each answer should be 1–3 sentences — specific, concrete, no marketing speak. Copy this template into a file called `business_context.md`.

#	Question
1	One-sentence company description
2	Industry / niche
3	Time in business
4	Customer transformation — BEFORE
5	Customer transformation — DURING
6	Customer transformation — AFTER
7	Customer transformation — PROOF
8	Core offer (name, price, deliverable, time-to-value)
9	Business model
10	Why you (and not a competitor)
11	Voice and forbidden phrases (5 words you say a lot, 5 you would never say)
12	Strategic narrative for this quarter (the single story all content should ladder up to)

A concrete answer set, for reference

1. We help solo HVAC contractors in Canada land 5–10 net-new commercial accounts per quarter without cold calling.
2. Vertical lead-gen for trades. Not general B2B marketing.
3. 4 years.

4. Before: contractor relies on word-of-mouth + Google Ads. Pipeline is whatever showed up.
5. During: we run their LinkedIn outbound + a property-manager intent list, weekly cadence.
6. After: 90 days in, 5-10 booked discovery calls per month from named target accounts.
7. Proof: 14 contractors, \$2.4M in net-new commercial revenue tracked through year 1.
8. "HVAC Pipeline Engine" — done-for-you outbound, \$3,500/mo + setup.
9. Subscription, 6-month minimum. Cash up front for setup.
10. We refuse to take more than 3 contractors per metro. Geographic exclusivity.
11. Say a lot: pipeline, named-account, dispatch, route density.
Never say: synergy, leverage, journey, ecosystem.
12. This quarter: "You don't need more leads, you need fewer better ones."

Exercise — your first context file

1. Create a file called `business_context.md` inside `01_Business_Context/`.
2. Block out 90 minutes. Disable Slack, email, and your phone.
3. Answer all 12 questions. If an answer is more than 3 sentences, you are explaining — cut.
4. Save the file. Attach it to every AI run from now on. No exceptions.

PITFALL

If you find yourself writing aspirational answers ("we will be the leader in..."), stop. The model writes from the present tense. Aspirations create slop.

Try it now — a 15-minute experiment

Before you read further or upgrade or do anything else: prove the principle to yourself.

1. Open a new chat with Claude or ChatGPT.
2. Run A: ask for a tweet on a topic you care about. Save the result.
3. Run B: paste in 5 sentences describing your business. Re-ask. Save.
4. Run C: paste in 20 of your best past posts. Re-ask. Save.
5. Read the three side by side. Where did the model start sounding like you?

WIN PATTERN

The first time you watch the same model produce three obviously different results — same prompt, different attached context — the framework clicks. It will not feel theoretical anymore.

What’s in the full course

This guide gave you the foundation: the principle, the folder structure, and the most important file. The full Practice Manual gives you the rest of the system — the prompts, the workflows, and the code.

Inside the 56-page Practice Manual

Part	What you learn
1. Foundation (you have part of this)	Why AI output reflects input quality. The 4-folder system. The 12-question business context. The anatomy of a real prompt.
2. Content workflows	The tweet generation loop. Long-form (YouTube + blog) scripts. Short-form (Reels / Shorts / TikTok). The repurposing engine: 1 idea → 6 outputs.
3. Vibe coding fundamentals	Setting up Claude Code or Antigravity. How web apps actually work (for non-coders). The vibe-coding development loop. Project 1: deploy a portfolio site in 30 minutes.
4. Building real SaaS	Project 2: a full-stack client dashboard with auth and real-time. Project 3: a lead-gen SaaS with payments. Project 4: a thumbnail-generator SaaS.
5. Production readiness	The 80/20 of security for vibe coders. Deployment, monitoring, iteration. Pricing, marketing, and launch.

Inside the Practice Kit

- Pre-built folder structure (just download and rename).
- Filled-in business_context.md template with prompts in every field.
- 6 production-grade prompt templates: tweet, long-form, short-form, vibe-coding rules, security audit, idea-to-spec.
- 3 starter code repos (portfolio, dashboard, lead-gen SaaS) with prompts ready to drop into Claude Code.
- Cheat sheets: prompt engineering levers, the default vibe-coding stack, glossary.

Three tiers

Tier	Price	What you get
Starter	\$79	56-page Practice Manual (PDF + DOCX). Practice Kit folder. 6 prompt

		templates. Lifetime updates.
Pro	\$249	Everything in Starter. 3 starter code repos. 8 video walkthroughs of the workflow loops. Private community.
Cohort	\$999 (cap 25)	Everything in Pro. 4-week live cohort. Weekly office hours. 1 hot-seat review. Certificate.

KEY IDEA

If you implemented the 15-minute experiment and the framework clicked, the Pro tier will pay for itself the first week you publish on-voice content or ship a starter project.

Get the full course

→ aimastery.example.com/course

Use code STARTER25 at checkout for \$25 off any tier this week.

Questions? Hit reply on any of the emails I'll be sending you over the next few days. I read every one.